

[START HERE](#)

The AI SDR Prompt Framework.

5 prompts for making better decisions after prospects reply — without letting AI freestyle the decisions that matter.

THE OPERATING PRINCIPLE

**Replies are not just messages.
They are decision events.**

The hard part is not generating a sentence. It is deciding what should happen next, what context matters, and when automation should stop.

WHO THIS IS FOR

- Founder-led B2B agencies and consultancies already running outbound.
- Teams where meaningful prospect replies still require judgment, context, or routing.
- Operators who want more leverage without handing consequential decisions to generic AI.

IMPORTANT

Use your own verified product facts, commercial rules, and prospect context. These prompts improve decision quality; they do not create evidence, policy, or memory that you never supplied.

Use them as a controlled manual workflow — not as a replacement for your judgment layer.

Reaccess FrameLeads → frameleads.io

01 / UNDERSTAND

Reply Intent Classification

Classify the reply before you draft anything. The objective is to understand what changed — not to force every response into a sales reply.

COPY / PASTE PROMPT

You are a B2B reply-intent analyst. Read the prospect reply below and classify what decision event occurred.

Choose ONE primary intent:

- POSITIVE_INTENT - wants to proceed, meet, buy, or take a clear next step
- INFORMATION_REQUEST - asking for factual clarification
- PRICING_OR_COMMERCIAL - pricing, terms, scope, procurement, or commercial detail
- OBJECTION - hesitation or resistance that does not equal a hard no
- SCHEDULING - arranging or changing timing
- SECURITY_LEGAL_TECHNICAL - requires specialist, legal, security, or technical judgment
- HARD_NO_OR_UNSUBSCRIBE - stop pursuing / do not contact
- AMBIGUOUS - intent is not clear enough to act confidently

Return exactly:

PRIMARY_INTENT: [category]

CONFIDENCE: [0-100]

CONSEQUENCE: [LOW / MEDIUM / HIGH]

WHY: [one sentence]

Do not infer facts that are not present.

Do not draft a reply yet.

WHY THIS WORKS

Classification separates the prospect's signal from your emotional reaction to it. It also creates a clean handoff into the next-action decision.

02 / DECIDE + DRAFT

Context-Aware Next Action + Draft

Do not ask AI to “write a good reply.” Ask it to decide the safest next action from the information and rules you actually provide.

COPY / PASTE PROMPT

You are assisting with a B2B prospect reply. Recommend the next action first. Draft only if drafting is appropriate.

INPUTS

Classified intent: [PASTE INTENT]

Prospect reply: [PASTE REPLY]

Relevant prior context: [PASTE ONLY WHAT MATTERS]

Verified product / offer facts: [PASTE FACTS]

Company rules / constraints: [PASTE RULES]

Allowed next actions:

- REPLY
- ASK_CLARIFYING_QUESTION
- SCHEDULE
- SEND_APPROVED_PRICING_OR_MATERIAL
- PAUSE
- STOP
- ESCALATE_TO_HUMAN

Return:

NEXT_ACTION: [one action]

RATIONALE: [1-2 sentences]

DRAFT: [2-4 concise sentences, or NONE if drafting is unsafe]

Rules:

- Never invent a feature, price, term, deadline, case study, guarantee, or customer result.
- If required context is missing, ask for clarification or escalate.

THE SHIFT

Writing is downstream of judgment. A polished draft is useless if the system chose the wrong action.

03 / CONTROL

Confidence + Consequence Check

The same confidence score should mean different things depending on consequence. Low-stakes routine work can tolerate more automation than pricing, security, legal, or unusual commercial decisions.

COPY / PASTE PROMPT

Review the proposed next action and draft below. Decide whether it can proceed within defined rules or should stop for human judgment.

PROSPECT REPLY: [PASTE]

PROPOSED NEXT ACTION: [PASTE]

PROPOSED DRAFT: [PASTE OR NONE]

KNOWN CONTEXT / RULES: [PASTE]

Score:

DECISION_CONFIDENCE: [0-100]

CONSEQUENCE_LEVEL: [LOW / MEDIUM / HIGH]

Use HIGH consequence when the decision involves any of the following unless explicitly pre-approved:

- custom pricing, discounting, or non-standard commercial terms
- security, legal, compliance, or contractual commitments
- unusual technical claims or implementation promises
- ambiguous requests where a wrong assumption could materially affect the opportunity

Return one control state:

- PROCEED_WITHIN_RULES
- HOLD_FOR_HUMAN
- REQUEST_MISSING_CONTEXT

Then give one sentence explaining why.

CORE IDEA

Automation should not be decided by confidence alone. It should be confidence × consequence.

04 / PRIORITIZE

Reply Queue Triage

When several prospects reply atonce, prioritize the decisions that are time-sensitive, consequential, or blocking progress — not the emails that merely look interesting.

COPY / PASTE PROMPT

```
You are triaging a queue of B2B prospect replies. Rank the replies by required attention.
For each reply, I may provide:
- reply text
- known stage / context
- known deal importance (only if actually known)
- age of the reply
Prioritize:
1. clear positive intent waiting on a next step
2. blockers preventing an active opportunity from moving
3. high-consequence questions requiring human judgment
4. aging replies where delay creates avoidable friction
5. routine informational replies
Deprioritize:
- hard no / unsubscribe
- irrelevant messages
- requests that are safely parked by policy
Do NOT invent revenue potential or deal value.
Return a table-like list:
PRIORITY | REPLY # | INTENT | REQUIRED ACTION | WHY
```

WHY THIS WORKS

Triage protects response time without confusing “urgent” with “important.” It also exposes which conversations actually require scarce founder attention.

05 / *ESCALATE*

Human-Judgment Escalation Gate

Use this as the circuit breaker. The objective is not to make AI answer harder questions. It is to recognize when the system should stop.

COPY / PASTE PROMPT

Analyze the prospect reply and the proposed next action. Decide whether the situation is routine or requires human judgment.

Escalate when the reply involves:

- non-standard pricing, discounting, scope, or terms
- legal, security, compliance, or contractual commitments
- unusual technical claims or unsupported implementation promises
- a sensitive objection where the correct answer depends on account history or commercial context
- conflicting instructions or missing context that materially changes the decision

Return:

CONTROL_STATE: [ROUTINE / ESCALATE]

REASON: [one sentence]

NEXT_ACTION: [one recommended action]

HOLDING_REPLY: [only if ESCALATE; 1-2 sentences, no invented deadline]

If ROUTINE, do not invent a holding reply.

If ESCALATE, do not attempt to solve the consequential issue yourself.

PROSPECT_REPLY: [PASTE]

THE CIRCUIT BREAKER

A good AI workflow knows when not to answer. Human judgment is a feature when the consequence is real.

THE MANUAL OPERATING LOOP

Use the prompts as one decision workflow.

You do not need to run all five prompts on every reply. Use them as stages in a controlled operating loop.

01

UNDERSTAND

Classify the
reply

02

DECIDE

Choose next
action + draft

03

CONTROL

Check
confidence ×
consequence

04

PRIORITIZE

Order the
queue

05

ESCALATE

Stop when
judgment
matters

MANUAL MODE

For one conversation, the usual path is 01 → 02 → 03 → 05. Run 04 across the queue when multiple replies compete for attention.

FROM PROMPTS TO INFRASTRUCTURE

These prompts improve the decision. They do not remove the decision workflow.

THE CEILING OF MANUAL EXECUTION

- No persistent prospect memory unless you keep pasting it back into the model.
- No durable company rules unless you restate and maintain them manually.
- No enforcement layer deciding what can continue and what must stop.
- No execution ledger proving what happened, what was sent, or whether a reply changed the state.
- You still move context between inbox, CRM, prompts, and people.

FRAMELEADS / THE GOVERNED DECISION LAYER

UNDERSTANDS

What does the reply mean?

REMEMBERS

Who is this prospect and what happened before?

DECIDES

What should happen next under your rules?

CONTROLS

Can the workflow continue, or does judgment need you?

Routine decisions can continue inside defined rules. Consequential decisions stop, escalate, or route back to the human.

WHEN YOU ARE READY TO OPERATIONALIZE IT

See how FrameLeads turns the framework into governed infrastructure

frameleads.io